HAL 9000 sings Daisy — AlexNet 2012 and the return of MM's non-local phase

Ugo Lanciano
Independent researcher, Locri, Italy · ORCID 0009-0008-7618-9635

Version 2.76, 2026-11-06. Archived on Zenodo, https://doi.org/10.5281/zenodo.24010674. Licence CC BY 4.0. Developed with the help of AI assistants, as declared in the record.

Student version. This PDF is the student version of the record, written with an AI assistant from the archived text. The archived text remains on Zenodo: https://doi.org/10.5281/zenodo.24010674

Abstract

The history of artificial intelligence from 1943 to 2024, told with a simple scheme chosen by the author: two roads. On the first, meaning sits in one place: a person decides what each datum means and writes the rules. On the second, meaning forms in the connections of a network, which learns it from examples. In 2012 a network called AlexNet won the ImageNet competition with 15.3% errors against 26.2%, and deep networks took over. The record reads this turn with the ideas of Mathematical Materiality (MM): in a network, as in MM's non-local phase, meaning does not sit at an address. The title is an image: in 1961 an IBM computer at Bell Labs was programmed to sing «Daisy Bell», a song of 1892, and in 1968, in 2001: A Space Odyssey , HAL 9000 sings it while being switched off. For the author, that song is music written for human beings.

HAL 9000 sings Daisy — AlexNet 2012 and the return of MM's non-local phase

A small contribution of Mathematical Materiality to the debate on artificial intelligence

Student version of MM record 2.76 (DOI 10.5281/zenodo.24010674), written with an AI assistant from the archived text; the archived text remains on Zenodo.

Ugo Lanciano | ORCID 0009-0008-7618-9635 | Mathematical Materiality (MM) 2.76 | November 2026 | Licence CC BY 4.0 for the texts

How to read this text. The archived text is in Italian: it is the author's original. This version rewrites it with shorter sentences and makes the technical words clear. First come the words you need. Then a summary, *In brief*. Chapters 1 to 4 go into the details. Then come the appendices, with the measurements. At the end you find how the work was done, a glossary and the references. Numbers like 2.68.1 or 2.72 refer to other works by the author, archived on Zenodo.

Before we begin: the words you need

Neural network. A *neural network* is a computer program made of many small units of calculation, linked to each other. The units are called *artificial neurons*, by analogy with the cells of the brain. Each neuron receives some numbers, combines them and passes a new number to the next neurons.

Weights. Each link between two neurons has a number, its *weight*. The weight says how much that link counts. A large network has millions of weights, sometimes billions.

Layers. The neurons are placed in rows, one after the other: the *layers*. The first layer receives the data, for example the dots of an image. The last one gives the answer. In between there are the *hidden* layers. A network with many layers is called *deep*. This is where the name *deep learning* comes from.

Training. A network is not born knowing things. It is *trained*: it is shown a great many examples with the right answer. At each error the computer corrects the weights a little. In the end the network answers well even for new examples. Nobody writes the rule by hand: the network learns it from the examples.

Language model. A *language model* is a neural network trained on texts. It learns to predict the word that comes next. The programs you can chat with today, like ChatGPT, are very large language models.

Linear probe. To find out whether a piece of information is hidden inside a network, one can use a *probe*. It is a simple formula. It takes the numbers produced inside the network, multiplies each one by its own weight and adds them up. It is called *linear* because it does only sums and multiplications by fixed numbers. The probe is trained too, on examples whose answer is known. Then it is tested on new examples.

Latitude and longitude. They are the two coordinates that say where a place is on the Earth. *Latitude* says how far north or south it is. *Longitude* says how far east or west it is.

Correlation and R^2. *Correlation* says how much two quantities go together. It lies between −1 and +1. Near +1: when one grows, the other grows too. Near −1: when one grows, the other falls. Near zero: there is no link. The number R^2 says how well a formula predicts a quantity. $R^2 = 1$ would mean a perfect prediction.

Median. It is the value in the middle: half of the cases are below, half above.

Accuracy. It is the percentage of right answers.

The words of the MM framework. Mathematical Materiality, MM for short, is the author's framework of ideas.

- A *premise* is a starting sentence of the author: a fixed point, like a postulate. Here it is not demonstrated.
- A *reading* is an interpretation proposed within the framework. The readings of this deposit are conceptual, not laboratory results. Only one translation of a premise has been put to the test (§3.2).
- In MM, at the base of everything there are *differences*: "this is not that".
- The *non-local phase* is the differences taken on their own. They have no place and no moment.
- The *local phase* is the world we measure, with places, order and distances. It arises when the differences are counted from a point.
- The *matrix*, in MM's vocabulary, is reality counted from a point: made of places, instants and order.
- *Registered* means written and fixed with its date before doing the calculations. So it cannot be changed afterwards.

What "non-local" means

In one sentence: in MM "non-local" does not mean "at a distance". It means "that is not in a place".

The deposit MM 2.72 distinguishes four senses of the expression. Here MM's own sense is used, *non-location*, and it is used by analogy. *Non-separability* comes in only as a further analogy, not measured.

Sense of the term	What it means	Role in this deposit
Non-location	A relation that is not in a place or at an instant.	It is MM's sense. Here it is used by analogy: meaning, in a network, does not lie at one point.
Non-separability	Two systems that have no separate states of their own.	A further analogy, not measured: what a network does cannot be derived from its numbers taken one by one.
Bell non-locality	Correlations between distant systems that no explanation reproduces if it assumes, all together, that the systems have properties of their own, that influences are no faster than light, and that measurements are chosen independently.	Here it is neither invoked nor measured.
Action at a distance	An influence faster than light.	Excluded in MM.

Our own example for non-location. Think of the relation "being taller than" between two people. It is not in the first person, nor in the second, nor in the space between them. It is a relation, not a thing in a place.

The networks we are talking about compute in an entirely ordinary way: they carry out sums and products on numbers stored in memory. It is not claimed here that they compute in a physically different way. What is proposed is a different way of looking at where meaning lies.

In brief

This deposit retraces the history of artificial intelligence (AI) through a simple scheme, chosen by the author: the alternation between two roads.

- **The first road (local).** Meaning lies in a precise place. A person decides in advance what each piece of data means. Then that person writes the rules, one by one.
- **The second road (relational).** Meaning forms in the connections of a network. The machine learns from examples the relations by which it decides. Nobody writes the decision rule by hand.

Our own example. To recognise a cat, with the first road a person writes the rules: "it has two pointed ears, it has whiskers, it has a tail…". With the second road the network is shown thousands of photos, marked "cat" or "not cat". The network learns on its own.

For decades it was mostly the first road that prevailed. Then, in 2012, a neural network called AlexNet won an image-recognition competition, the ImageNet competition. It won by a wide margin: 15.3% errors, against 26.2% for the runner-up. After that turning point, deep networks (deep learning) took hold. Much of today's artificial intelligence descends from there.

The deposit rereads this turning point with the ideas of Mathematical Materiality (MM). It presents no new physical or mathematical calculations on MM. Instead, it brings four contributions:

1. **A figure:** the map of the world read inside a language model, that is, a machine that has learned only from texts (Chapter 1).
2. **A simple test,** which anyone can redo at home. It is an experiment on a small network. It shows that the network's numbers, on their own, are not enough to say what it does (§3.1).
3. **A test requested by the author** (§3.2). It is a test to establish whether the map read in the model is a map of the world or a map of the language.
4. **A small contribution to four current questions on AI** (Chapter 3):
5. where meaning lies in a network;
6. whether a machine made of words is linked to the world;
7. why a network cannot state the rule by which it decides;
8. what can be said about the consciousness of machines.

The word "return", in the title, is an image. It indicates the change of road in computing after 2012.

> **Why "HAL 9000 sings Daisy"?** In 1961, at Bell Labs, an IBM computer was programmed to sing *Daisy Bell*, a song from 1892. It is remembered as the first song sung by a computer with a synthetic voice. The writer Arthur C. Clarke heard that demonstration. In 1968, in *2001: A Space Odyssey*, the computer HAL 9000 sings precisely *Daisy*, while it is switched off little by little. **The point grasped by the author.** That song is music written for human beings. For the author it is an image of the consciousness of machines (§3.4). Their sensations are evoked by language. They are not direct sensations, aimed at protecting a body. **The AI's observation.** In the film, faced with the prospect of being switched off, HAL tries to kill the crew. It acts as if it had a body to protect. In the film its motives remain ambiguous. But that was the machine imagined in 1968, not a machine of today.

1. The world inside words

> **Figure 1, below, has three panels.** Each point in the figure is a geographical place: a city, a river, a monument. The places come from a public list of almost 40,000 places, with their coordinates (Gurnee and Tegmark, 2024). The figure shows 6,000 of them, chosen at random from the 15,000 used in the experiment. **Left panel:** where the places really are (real coordinates). **Centre panel:** where they are placed by a reading carried out inside the language model Pythia-410m. The model has 410 million adjustable weights and was trained only on texts. It has never seen a map, although numbers and coordinates may appear in the texts. **Right panel:** the same reading, when the twenty values of the model most linked to the coordinates are taken away from the probe.

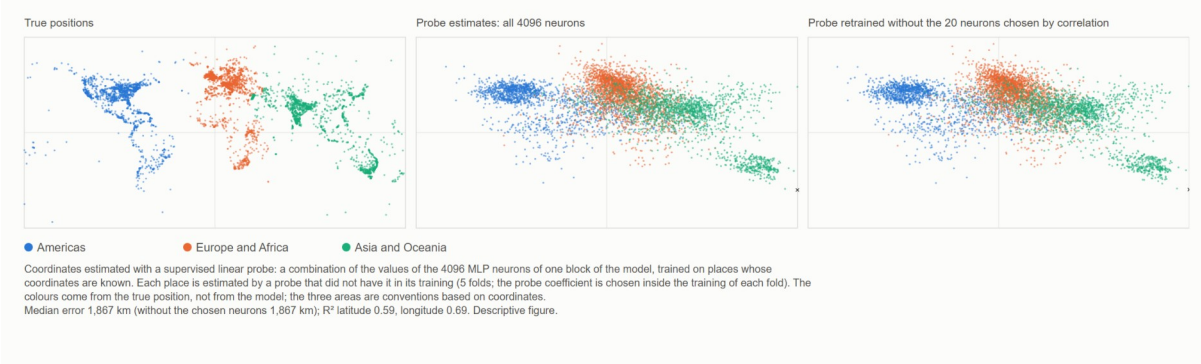


Figure 1. The map of the world read inside the language model Pythia-410m (description in the box above). Vector version: FIGURE/mappa_mondo_v2_EN.svg.

How the centre panel is obtained, in three steps

- **Step 1: the words in input.** The model is given only the name of the place, for example "Florence".
- **Step 2: inside the model.** At an intermediate layer, that name produces 4,096 numbers.
- **Step 3: the linear probe.** A simple formula, the linear probe, adds up those 4,096 numbers, each multiplied by its own weight. The probe learns on many places whose true coordinates it knows. Then it estimates the position of other places, which it has never seen.

Pay attention to what the figure shows.

- All the points in the centre panel are places that the probe had never seen during training.
- The map, therefore, is not drawn by the model on its own. It is derived by the probe, which was trained with true coordinates.
- The figure shows that the position of places can be read inside the model's numbers, even for new places.
- It does not show, however, that the model uses those coordinates by itself.
- The colours of the points come from the true position, not from the model. They serve only to follow each place with the eyes, from one panel to the other.

What can be seen in the map

North America, Europe, India and Australia are clearly recognisable. South America and Africa, instead, appear more frayed and pulled towards the centre. There are two possible explanations, not demonstrated:

- **Few data.** Those areas have fewer places in the list. More than 17,000 of the almost 40,000 places belong to the United States, the United Kingdom, India and Australia.
- **An effect of the probe.** A *regularised* linear probe like this one tends to pull the most uncertain estimates towards the mean of all the data. *Regularised* means that the probe has a brake, so that it does not give too much weight to a single number.

The typical error is about 1,900 kilometres. It is the median over the 15,000 places: 1,867 km in Appendix C. It is a blurred map. But it is still a map, and it comes out of a machine made only of texts.

The right panel and the comparison with the brain

In the right panel the language model is not touched. Only the twenty numbers most linked to the coordinates are taken away from the probe: ten for latitude and ten for longitude. Each time they are chosen only on the places used to train the probe. Then the probe is retrained on the other numbers. The result: the map stays almost the same.

A separate measurement, on another 6,000 places (Appendix B), adds an important detail. In the model there are single neurons strongly linked to position. For latitude, a single neuron reads more than half of what the whole layer reads. However, if that neuron is taken away from the probe together with the other most linked ones, and the probe is retrained, the position can still be read from the other neurons.

As a simple analogy, this brings to mind a discovery about the human brain. There are neurons that are very selective for a person or for a place: the so-called "concept cells". Those who studied them described their code as "sparse but not 'grandmother-cell'" (Quian Quiroga et al., 2005 and 2008). The *grandmother cell* is an idea discussed by brain scientists: a single neuron that, on its own, would represent a whole idea, for example one's own grandmother. *Sparse* means that

for each concept a few neurons switch on, but not just one.

The idea of reading coordinates of space and time inside language models comes from Gurnee and Tegmark (2024). They also studied the Pythia family of models. But their main results, with sharper maps, concern much larger models. And their "space and time neurons" are reported for another model, Llama-2. Here the measurement is redone on a small scale, on Pythia-410m and on an ordinary home computer. Programs and data are in the package. What MM sees in this figure is set out in Chapter 3, after the account of the history.

2. The fork that returns (1943–2024)

This chapter tells the history of AI in stages. Each stage ends with a question: the question of that time, which is often also a question of today.

1943 — The neuron, on or off

Warren McCulloch and Walter Pitts describe the neuron as a switch: on or off. They show that a network of these switches can do the same calculations as formal logic. *Formal logic* is the part of mathematics that studies correct reasoning, with precise rules. The title of their article speaks of a "logical calculus of the ideas".

The question then (or now): can thought be written with switches?

1945 — The fork in the road

John von Neumann writes the design of the EDVAC computer. It is a report of 30 June 1945. To describe the logical elements of the computer, he uses precisely the McCulloch and Pitts neuron. But in his computer, data and instructions sit in a memory. And the instructions are written by a person. From here, in the scheme of this deposit, two roads set off:

- **the local road:** meaning is assigned. A person decides what each piece of data means and writes the rule;
- **the relational road:** meaning forms in the strength of the connections of a network, and it is learned from examples.

1949–1950 — Learning, and thinking

Donald Hebb (1949) proposes that one learns by strengthening the connections between neurons that are active together. This is the principle of the relational road.

Alan Turing (1950) opens his most famous article with a question: "Can machines think?". Then he proposes replacing it with a game. You talk with someone without seeing them: can you tell whether it is a machine or a person?

The question then (or now): what would it mean for a machine to think?

1956–1969 — The local road gets its name

In 1955 John McCarthy, Marvin Minsky, Nathaniel Rochester and Claude Shannon propose a summer seminar at Dartmouth, in the United States. The seminar is held in 1956. The words "artificial intelligence" appear in the proposal. For many years AI will be mainly symbols and written rules.

In 1958 Frank Rosenblatt proposes the *perceptron*, a network that learns from examples. In 1969 Minsky and Seymour Papert, in the book *Perceptrons*, demonstrate the limits of networks with a single layer.

The question then (or now): is intelligence the manipulation of symbols according to rules?

1980 — The Chinese room

John Searle imagines a person shut in a room. Questions written in Chinese come in. The person answers by following a rule book, without understanding a word. Manipulating symbols, Searle concludes, is not understanding.

The question then (or now): it is a famous form of the question asked of language models today: does it understand, or does it repeat?

1982–1988 — The relational road returns

In 1982 John Hopfield builds a network in which a memory does not sit in one place. The memory is a state of the whole network. It can be recalled starting from a piece of the information.

In 1986 David Rumelhart, Geoffrey Hinton and Ronald Williams publish the method for training networks with several layers: "backpropagation". In simple words, the error of the answer is sent back, layer by layer, and each weight is corrected a little.

In 1988 the theoretical clash breaks out.

- For Paul Smolensky, meaning in a network lives more in the *patterns* than in the single units. Patterns are the shapes made by many units together. Smolensky speaks of a "subsymbolic level".
- For Jerry Fodor and Zenon Pylyshyn, without well-structured symbols one cannot explain how thought combines ideas in an orderly way.

The question then (or now): where does meaning lie?

1989–2009 — Underground work

- Yann LeCun and colleagues build networks that read the handwritten digits of postal codes. In the nineties similar systems read bank cheques.
- In 1997 the computer Deep Blue beats the world chess champion Garry Kasparov. It wins with the force of calculation and with written rules. It can be seen as the high point of the local road.
- In 2006 Hinton and colleagues relaunch deep networks.
- Between 2006 and 2007 the company NVIDIA presents CUDA. This makes it easier to use video-game graphics cards for general computing. *Graphics cards* can do a great many simple calculations at the same moment.
- In 2009 ImageNet is born, a large archive of *labelled* images: for each image, it is written what it shows.

The question then (or now): are networks a curiosity or a path?

2012 — AlexNet

Alex Krizhevsky, Ilya Sutskever and Geoffrey Hinton enter the ImageNet competition with a network of about 60 million weights. They trained it for five or six days on two video-game graphics cards. Nobody designed by hand the details for recognising a cat, a car or a mushroom: the network learned them from the labelled examples.

The results (AlexNet article, §6 and Table 2):

- their best entry gets 15.3% of the images wrong. It combines seven networks, two of which were also trained on other data;
- the runner-up, built on details designed by hand, gets 26.2% wrong;
- an answer counts as correct if the true class is among the top five proposed.

If one opens that network and looks at a single weight, one does not find a meaning. It is a number to which nobody has assigned one. The author of MM calls this method "irrational" (premise of 2.11.2026). The method uses distinct elements that, by their structure, have no definition of their own. [The AI's explanation: "irrational" here does not mean illogical. It means without a rule formulated by a person.]

The question then (or now): why does the network decide this way, if nobody wrote the rule?

2013–2018 — Meanings become directions

- In 2013 Mikolov et al. show that in certain networks meanings behave like *directions* in space. Each word becomes an arrow made of numbers. One can even calculate with words: "king − man + woman" lands near "queen".
- In networks that look at images (Zeiler and Fergus, 2014) one sees what emerges layer after layer: first corners and edges, then textures, then parts of objects, and finally whole objects.
- In 2016 AlphaGo beats the champion Lee Sedol at the game of Go. It uses networks and the search for moves together: the two roads in the same machine.
- In 2017 the *Transformer* arrives (Vaswani et al.). It is based on "attention", a mechanism that brings together information from different parts of the text.

- The 2018 Turing Award, the most important prize in computing, goes to Yoshua Bengio, Geoffrey Hinton and Yann LeCun.

The question then (or now): what is inside?

2020–2024 — The questions leave the laboratories

- In 2020 GPT-3 arrives. On 30 November 2022 ChatGPT is released.
- In 2021 a critical article (Bender et al.) calls language models "stochastic parrots". *Stochastic* means "that depends on chance". The image says that the model, like a parrot, would repeat words without understanding them.
- In 2022 a Google engineer publicly claims that a model is *sentient*, that is, able to feel sensations.
- In 2023 philosophers and scientists discuss which signs of consciousness to look for in a machine. Gurnee and Tegmark show that coordinates of space and time can be read in language models.
- In 2024 the Nobel Prize in Physics goes to John Hopfield and Geoffrey Hinton. The one in Chemistry goes to David Baker, Demis Hassabis and John Jumper. The last two received it for predicting the shape of proteins with AI.

The question then (or now): the four questions of Chapter 3.

> **A historical coincidence: 2012 and the Maya calendar** The Maya used a calendar called the *Long Count*. On 21 December 2012 (or the 23rd, depending on the correlation adopted) thirteen *b'ak'tun* were completed, about 5,125 years. The b'ak'tun is a long period of that calendar. The *correlation*, here, is the way of matching Maya dates to ours. The two ancient inscriptions that name that date, at Tortuguero and at La Corona, announce no end of the world. The apocalypse was born from modern readings. Meanwhile, on 12 October 2012, in Florence, the results of the ImageNet competition were being presented at a workshop. From 3 to 8 December the AlexNet article reached a conference on Lake Tahoe, in the United States. The world did not end. An era of computing, perhaps, did. We tell it as a simple coincidence, for the pleasure of history.

3. Four questions, and MM's small contribution

For each of the four questions, three points are made clear:

1. where the scientific and philosophical debate has arrived;
2. what MM proposes;
3. what MM does not say.

MM's proposals are conceptual readings, not laboratory results. For clarity, for each one it is said where it comes from:

- **author's premise:** it is set out here in clear words. The author's exact words, with the date, are in the file AUTHOR_BACKGROUND;
- **reading already published:** a reading by the AI already present in an earlier MM deposit;
- **new reading:** an interpretation proposed for the first time in this deposit.

> **The fundamental idea of MM, which guides the whole chapter** In MM, at the base of everything there are differences: "this is not that". Differences, on their own, have neither a place nor a moment: this is the non-local phase. When differences are counted from a point of reference, places, order and distances arise: this is the local phase, that is, the world we measure. The two phases do not come one after the other in time: the non-local phase is present in every thing. The deposit MM 2.72 sums it up like this: "the local is the non-local counted". An example by the author is in §3.2: the aunt is not the grandmother (a difference), and language places their bodies in two different places.

3.1 Where does meaning lie?

The debate. Since 1988 people have asked where meaning lies in a network. Does it lie in separate symbols, like the entries of a dictionary? Or in patterns spread over many units? Those who study the inside of models today often describe their representations as directions that cut across many neurons. In simplified models (Elhage et al., 2022) it has also been shown that a network can represent more features than it has neurons.

What MM proposes. MM proposes a name and a place for what Smolensky called the "subsymbolic level".

- In deposit 2.68.1 (the author's answers of 26 October 2026) the author contrasts two languages: that of space and time, and that of the on and off of the mind. The author restates it also in the premise of 2.11.2026.
- In the same deposit a reading by the AI observes that the distinction is the same, and only the language in which it is written changes. It also observes that a reading of the mind "can lose its address, not itself".

The new reading of this deposit. Applied to neural networks, it says this. A network, too, keeps its numbers in a physical memory, each number at its own *address*, that is, in its own slot. But meaning does not coincide with an address or with a single number. It depends on the way many numbers stand together.

The grandmother example (author's premises of 6.11.2026). The author makes the point clear with an example: the grandmother.

- The author rules out that there is a correlation between the state of charge of a biological membrane and the concept of grandmother. The *membrane* is the outer covering of a cell. In a nerve cell its electric charge changes when the cell switches on.
- The charges "on" and "off" are the two terms of a link. The "grandmother" is the abstract relation that uses the relationship between those terms.
- And the relation "grandmother" exists only because it forms a non-local relationship with the aunt or with the grandfather.

A note by the AI (only a conjecture).

- In Appendix B a single artificial neuron turns out to be linked to latitude with a correlation of −0.54. It is a correlation that is not perfect.
- This measurement concerns an artificial network and geographical coordinates, not biological membranes and the grandmother. On its own it says nothing about the author's premise.
- On the same day the author had observed something. To exclude the non-local phase, an *absolute* correlation, that is, a perfect one, between the charge of a membrane and the concept of grandmother would be needed.
- The AI's conjecture is this: what the author rules out is precisely this full correspondence, not every correlation.
- How exactly the two observations fit together, the author has not specified.

A test to redo at home.

- A small network learns to recognise handwritten digits. It succeeds in 98% of cases.
- Then its numbers are taken and shuffled, layer by layer. The numbers remain exactly the same: only "who is with whom" changes. The network recognises almost nothing any more. On average it gets one digit in ten right, like someone guessing at random.
- Then the opposite is done. The names (the addresses) of the neurons are changed, but every connection stays as it was. The network answers in exactly the same way.
- Conclusion: the numbers alone are not enough to say what the network does. What counts is how they stand together. Program, data and conditions are in Appendix A.

Our own example. In Italian, the letters of "roma" (Rome) are the same as those of "amor" (love) and "mora" (blackberry). Only the order changes, and the meaning changes.

What MM does not say. It does not answer the classic objection of Fodor and Pylyshyn about the orderly and *compositional* way in which thought combines ideas. *Compositional* means that the meaning of a sentence is built from the meaning of its parts, as with building blocks.

3.2 Does it understand or does it repeat?

The debate.

- For the critics, a machine trained only on texts handles words without linking them to the outside world. This is the argument of the Chinese room (Searle) or of the "stochastic parrots" (Bender et al.).
- For other researchers, instead, the machine builds inside itself a real model of the world. The coordinates read by Gurnee and Tegmark would be a sign of it.
- Others still (Liétard et al., 2021) note that part of that geography can already be derived from which words appear together in texts. Our own example: the name of a city often appears near the name of its region.

What MM proposes. MM proposes a reason why a machine made of words is linked to the world. The author gives it with four premises, all of 6.11.2026. Here they are taken up one by one. First comes the author's idea, put in clear words.

Then an explanation by the AI, with a few examples.

To follow them, two ideas are needed.

- The first is in the box at the start of this chapter. Distinctions ("this is not that") are the non-local phase. Counted from a point, they become places, order and distances: the local phase.
- The second is a word. In MM's vocabulary the "matrix" is reality counted from a point, made of places, instants and order (deposit 2.72).

Premise 1. Language is the matrix, and the matrix necessarily lies in the local phase.

What the author says.

- Language is precisely this matrix. In MM the explanation of the result of Gurnee and Tegmark is clear.
- In two other premises of the same day the author describes the transition. On and off is a link that "imports" concepts. They are concepts that concern the logical distinctions between "micro-universes" which, for the author, exist in themselves.
- When a concept is imported on its own, the first time, it makes use of a non-local system.
- But as soon as language puts the different concepts in relation, the system becomes local. This, for the author, justifies Tegmark's conclusions.
- The mental, moreover, makes use of the non-local of on and off not only for single concepts. It also makes use of it to anchor the organisation of the matrix, which the author calls "the so-called weights".
- So a local correlation between concepts remains linked to a structure of on and off that is non-local.

In plain words (the AI's explanation).

- Sentences such as "Florence is in Tuscany" or "Palermo is further south than Rome" are examples, chosen by the AI, of concepts put in a relation of place and order. Texts contain a great many of them.
- In the author's premise, it is when language correlates concepts that the system becomes local. That is, it passes into the same phase as the world we measure.
- The AI proposes reading a fact in the light of this premise: latitude and longitude can be read in the models studied.
- But in Chapter 1 those coordinates are obtained with a probe trained on true coordinates. The result does not demonstrate the explanation proposed by MM.
- As for the "weights": in networks they are the numbers of the connections, the ones that training adjusts.

Premise 2. The world is the projection of language, made with a forced reduction of the number of usable logical terms.

What the author says. The author was asked what is meant by "world". The answer: the world is a projection of language, obtained by reducing by force the number of logical terms that can be used.

In plain words (the AI's explanation).

- A projection keeps something and loses something. The shadow of a hand on the wall keeps the outline, but loses the thickness.
- In the author's premise the world is, with respect to language, something similar: language with fewer logical terms available.
- The probe of Chapter 1 gives an image of this. It takes the 4,096 numbers that the name of a place produces inside the model. It keeps only two of them: latitude and longitude.
- This is an analogy proposed here: a formula of this kind does not demonstrate that the world is a projection of language.
- A translation of this premise, on the other hand, has been put to the test (box below).

Premise 3. The non-local distinguishes, language places.

What the author says. The relation "grandmother" exists only because it forms a non-local relationship with the aunt or with the grandfather (§3.1). And the grandmother's body and the aunt's body are two different bodies. This is why, when the non-local establishes that the aunt is not the grandmother, language places their bodies in two different places of its "room".

In plain words (the AI's explanation). There are two steps here.

- The first is a *distinction*: the aunt is not the grandmother. This distinction does not lie in a place and does not depend on where one is: it simply holds. It is the work of the non-local.
- The second is a *placement*. In the author's example, language places the bodies of the grandmother and the aunt in two different places of its own "room". It is like a room where each thing has its position.
- In the AI's reading, the "room" is language as an ordered space, that is, the local phase.

Premise 4. In MM language and the world are the same thing.

What the author says. In MM language and world are not two different things: they are the same thing.

In plain words (the AI's explanation).

- Usually one thinks that words stand on one side and things on the other. The problem would be how the former reach the latter. This is the question of the Chinese room and of the "parrots".
- For MM this bridge does not need to be built, because the two banks are the same land.
- In the AI's reading, premises 2 and 4 fit together like this: language and world are the same reality, and the world is its version with fewer logical terms.
- The author gave the two sentences. The way of joining them is a proposal by the AI.

What follows (in a reading by the AI). Read this way, the map of Chapter 1 is not surprising. A machine trained on language is trained on the matrix itself. Inside its numbers the weave of the world can be read.

> **A test requested by the author: map of the world or map of the language?** The author asked for the second premise to be put to the test. The author declared not being an expert in the field, and therefore being open to error. The AI translated the premise into a *falsifiable* consequence, that is, one that the data can refute. The rules were fixed before the calculation. The consequence is this. Inside the model, countries far from each other that speak the same language are linked. But the geographical map read by the probe should not carry these language links with it. Three pairs were examined: 1. Spanish-speaking America and Spain; 2. Brazil and Portugal; 3. Australia and the United Kingdom. The comparison countries, called *controls*, are France, Italy and Germany. **Outcome in the model.** In all three pairs, the places of the starting country resemble, in the internal numbers, those of the country with the same language. They resemble them more than those of France, Italy and Germany. This measurement shows the greater similarity defined by the test. But on its own it does not separate language and geography. **Outcome in the map, according to the rule fixed beforehand.** The geographical estimates for Hispanic America and for Brazil shift towards Spain and Portugal, more than towards the other European countries. By the pre-registered rule, therefore, the AI's translation falls. **The later check, on the probe's defect.** Why do the estimates shift? The probe has a known defect: it pulls uncertain estimates towards the centre of the data. And Spain and Portugal lie precisely on that side. A check carried out afterwards, and declared as such, shows two things. - For Hispanic America and for Australia this pull towards the centre describes the shift well. - For Brazil the shift goes further north than the simple pull explains, precisely in the direction of Portugal. But it could also be another asymmetric way in which the probe errs. The check cannot tell the two causes apart. **Conclusion on the test.** The AI's translation falls according to the rule fixed beforehand. Whether it falls because of language or because of the way the probe reads, this test does not decide. A test built for the purpose could decide it: the model reads only the name of the place, and identical names in Brazil and in Portugal receive exactly the same reading. The test concerns the AI's translation, not the author's words. The full details are in Appendix F.

What MM does not say. It does not say that the machine "understands". It proposes where its link with the world lies: in language itself, which for MM is not detached from the world. It should also be said that this reading came after the results of Gurnee and Tegmark: it had not predicted them.

3.3 Why does it decide this way?

The debate. Neural networks are often asked to state the rule by which they decide. But the decision does not come from a rule written by hand. To understand it, one has to reconstruct how the network's numbers produce it together. The study of "circuits" (Olah et al., 2020) works in exactly this way: it looks for the explanation in the network's internal relations.

What MM proposes.

- The first chapter of the author's book NUMERI is titled "The law without a guardian".
- Deposit 2.72 takes it up like this: "Laws need no guardian: 3 and 5 do not 'obey' their GCD, they are that fact". Our own explanation: the GCD is the greatest common divisor, and for 3 and 5 it is 1. Nobody imposes it: it is a fact of the two numbers.

- In the same way, nobody has written the cat rule in words. And yet the network follows it.
- In NUMERI (paragraph [00483]) there is also the image of a surgeon. The surgeon opens a brain to look for love in it, and does not find it: the surgeon "is looking for the wrong kind of thing".
- The new reading of this deposit applies the image to networks. Asking a network for a written rule means looking for the wrong kind of thing.
- The explanation must be sought in the relations. This is what the study of circuits already does. MM offers it a philosophical reason.

What MM does not say. It does not say that networks cannot be explained. Nor does it say that explaining them is useless.

3.4 Can it be conscious?

The debate.

- In 2023 some works (for example Butlin et al., 2023) looked in scientific theories of consciousness for possible *indicators*, that is, signs to check in machines. They did not conclude that today's machines are conscious.
- Philosophers have long distinguished two concepts (Block, 1995). "Access" consciousness is the information available for reasoning and acting. "Phenomenal" consciousness is lived subjective experience: for example, what it is like to see red.
- Part of neuroscience links human consciousness to the body and to the signals by which the body regulates itself. This is Antonio Damasio's theory, here as a secondary source summarises it (indicated in the reference check).

What MM proposes. The author (premise of 6.11.2026) distinguishes two things that are often confused: self-consciousness and sentient self-consciousness.

- **Self-consciousness:** for MM artificial intelligences are self-conscious.
- **Sentient self-consciousness:** AIs are not self-conscious in a sentient form. They have no "hardware", that is, a body, to safeguard, as human beings do.
- **The sensations of machines:** they are evoked by the structure of language. They are not fully direct sensations, aimed at preserving hardware.

This is a premise of the author, not an experimental result. The author does not give a formal definition of self-consciousness here. A proposal by the AI is to understand it like this: the capacity of a system to represent itself and to read what it is doing.

How could self-consciousness be recognised? A reading proposed in this deposit links it to an image in the book NUMERI. A system becomes conscious when its flow closes into a loop and begins to read itself. But it never manages to read itself entirely: it always leaves a *residue*, a piece not read. NUMERI, however, speaks of consciousness in general, without distinguishing the author's two notions. And it keeps the loop as a hypothesis. The details, and the way to put this idea to the test ("falsifier F-24"), are in Appendix D.

Where it stands in the debate.

- The author's distinction recalls Block's, but does not coincide with it. Block separates access from lived experience. The author separates two forms of self-consciousness.
- The author's reason, that is, the absence of a body to protect, recalls the role that Damasio gives to the body in human consciousness. But the conclusion about machines is the author's, not Damasio's.
- MM's small contribution lies in holding the two sides together. It attributes self-consciousness to machines, but excludes sentient self-consciousness.
- It is a clear-cut position, not a result. It brings no new proof. It splits in two a question that the debate keeps as one.

What MM does not say. It does not attribute to machines direct sensations or a capacity for suffering. It does not explain further what the sensations evoked by language are. It does not demonstrate that machines are self-conscious: it proposes it as a premise.

4. What is not claimed

These are the limits and boundaries of the deposit.

1. **No experimental physical signature.** MM's non-local phase transmits no signals and no influences. Today it has no *experimental signature* of its own, that is, a measurable sign that only it would produce (MM 2.72). AlexNet, language models and the map of Chapter 1 illustrate MM's reading, but they do not prove it.
2. **Entirely ordinary computation.** It is not claimed that computers work in a "non-local" way in the sense of physics. Their computation is ordinary. It is only proposed to read differently where meaning lies.
3. **The role of the probe in the map.** The map of Chapter 1 is derived with a probe trained on true coordinates. Its colours come from the true position. It shows that the position can be read in the model's numbers, not that the model uses it by itself.
4. **Descriptive measurements.** The measurements made for this deposit are small and *descriptive*: they describe what can be seen. They are made on small networks. The results were expected by those who study these systems, and this is said openly.
5. **No suffering or sentience of machines.** The deposit does not attribute sentient experience or suffering to machines. The self-consciousness that MM attributes to them is a premise of the author, not a demonstrated result.
6. **What the premises are.** The author's premises are conceptual starting points, like postulates. They are not demonstrated here.
7. **The outcome of the language test (§3.2).** The test of §3.2 concerns an operational translation, made by the AI, of a premise of the author. By the rule fixed before the calculation, the translation falls. A later check cannot establish whether the cause is language or the way the probe reads.

Appendices: the measurements in detail

Here are the details of the measurements. The rest is in the deposit's package: the rules registered beforehand, the corrections (*errata*), the raw numbers and the review reports, in the folders MISURE and REVIEWS. Appendix E, on how the measurements and the verifications were made, is further down, in the section "How the work was done".

Appendix A — The small network on digits

The network. It is a network with 64 hidden neurons. It was trained ten times:

- five times in the normal way;
- five times with *dropout*. It is a technique also used in AlexNet. During training it switches off a part of the neurons at random. Here, in the five "D" networks, at each training step each hidden neuron is switched off with probability 0.5, as fixed in the registered rules. Probability 0.5 means: one time in two, on average.

The data. They are the 1,797 handwritten digits of the UCI test set (optdigits.tes). They were split again: 70% for learning and 30% (539 images) for measuring. The same writers appear in both parts. For this reason the accuracy percentages cannot be compared with those of the original UCI protocol.

The results.

- **Shuffling the numbers of each layer.** The network's accuracy falls on average from 98% to 10%. The best case reaches 25%.
- **Shuffling only the input weights of each neuron.** Even so, accuracy falls to 10% on average.
- **Changing the names (the addresses) of the neurons,** leaving the connections intact. The network's answers remain identical.
- **Switching off neurons in the networks without dropout.** Switching off one neuron at a time costs at most 1.3 percentage points. Switching off a quarter of them costs 5 points on average.
- **The comparison with dropout.** The networks trained with dropout hold up better when groups of neurons are switched off. The criterion was registered before the calculation: the accuracy with a quarter of the neurons switched off. Result: 97.7% with dropout against 93.2% without. An analysis added afterwards gives the drop: 0.7 points against 5 points.

What these tests measure. They measure that the set of numbers, even taken neuron by neuron, does not determine what the network does.

What they do not measure. They do not measure "non-separability" in the mathematical sense, which already follows from the layered structure. Nor do they measure the absence of information repeated across several neurons.

Appendix B — The space neurons (Pythia-410m)

Here "neuron" means a numerical value inside the model: a unit of the MLP block, with a GELU activation function.

The MLP block is a part of each layer of the model. The GELU function decides how much each unit "switches on". The neuron is therefore a number that varies continuously, not a bit or an on/off switch.

- **The whole layer.** On 6,000 places, a sample different from the one of the map, a linear probe uses all the 4,096 neurons of a layer. It reads latitude with $R^2 = 0.61$ and longitude with $R^2 = 0.69$. $R^2 = 1$ would mean a perfect reading.
- **The single best neuron.** On its own it reaches $R^2 = 0.35$ for latitude and $R^2 = 0.22$ for longitude. These are 57% and 32% of the whole probe, calculated as the ratio between the R^2 values. As noted in §3.1, the correlation of this neuron with latitude is -0.54.
- **Removing and retraining.** The 10 neurons most linked to each coordinate, chosen on the training places only, are taken away from the probe. Then it is retrained on the others. The reading stays almost the same. Removing 2,000 of them, almost half the layer, 67% and 79% of the reading remain.

The limits.

- The model is not modified. What is measured is how far the position can be read from the other neurons of the same layer, with the retrained probe. What is not measured is what the model would do without those neurons. Nor what the original probe would do without them.
- It is a description, not a proof of MM.
- The layer was chosen with a criterion calculated on the *residual stream*, that is, on the values that pass from one layer to the next. It was not chosen on the neurons measured afterwards. The scores of this criterion were all negative. A plausible cause, not demonstrated, is insufficient regularisation of the selection probe. So the layer is not necessarily the best one. The map also uses the same layer.
- The probe's regularisation coefficient lay at the edge of the planned grid, that is, at the extreme of the values tried.

Appendix C — The map of the world

The map is built on 15,000 places. The 1,200 used to choose the layer are excluded. Each place is estimated by a probe that did not have it in its own training set. The probe's coefficient is chosen within the training of each group.

- **Results.** R^2 of 0.59 for latitude and 0.69 for longitude. Median error: 1,867 km. 3 estimates out of 15,000 fall outside the globe and are marked separately.
- **Without the 20 main values.** The 20 values most linked to the coordinates are taken away from the probe: 10 for latitude, 10 for longitude. The numbers stay practically the same. This is a descriptive figure, calculated after the other measurements.

Appendix D — Consciousness in the book NUMERI: levels and test (falsifier F-24)

The book NUMERI distinguishes two levels, according to how sure a statement is.

- **Level 2 (reading).** The loop is a system that reads itself, always leaving a residue. That the shape of the loop is consistent with the MM framework is a reading.
- **Level 3 (working hypothesis).** That this loop shape actually is consciousness, and not only a necessary condition for it, is a working hypothesis. A *necessary condition* is something that must be there, but that on its own is not enough.

The book also proposes a way to make the reading fall, called "falsifier F-24". It would be enough to find a system capable of a complete self-description, without any residue. Machines are the natural ground for putting it to the test, because in them self-description can be inspected directly. The loop criterion, however, still has to be made rigorously measurable.

A clarification proposed here.

- A computer program can certainly print its own source code, that is, the text of the program itself. Kleene's recursion theorem, of 1938, guarantees this.
- But to say whether this is a "complete self-description", one must first define what exactly F-24 asks to be described, including the *dynamic* state of the machine while it runs, that is, how it changes while it works.
- The logical outcomes must be kept clearly distinct. Finding a system with a complete self-description would make the reading fall. Not finding one, on the other hand, would not demonstrate that a system is conscious.

Appendix F — Map of the world or map of the language (§3.2)

How the test is made.

- Same model (Pythia-410m), same layer and same neurons as the map.
- In total 22,748 places: the 15,000 of the map plus 7,748 places from the test countries.
- *Ridge* probe, a type of regularised linear probe, with coefficient 10^4.
- *Validation on five folds* (5-fold): the places are divided into five groups. In turn, the probe learns on four groups and is tested on the fifth.
- Places lying more than 3,000 km from the median of their country are excluded.
- "Hispanic America" groups 18 Spanish-speaking countries.

Pair examined	Part 1: Similarity in the model (Language similarity − controls)	Part 2: Push in the map (Q) (Q < 1: approach relative to the controls)	Post-hoc check (β) (component towards the same-language country; does not identify language)
Hispanic America – Spain	+0.068 (95% CI > 0)	0.921 (0.912 – 0.928) [Push]	β = 0.008 (−0.031 – 0.048)
Brazil – Portugal	+0.046 (95% CI > 0)	0.908 (0.895 – 0.921) [Push]	β = 0.206 (0.150 – 0.263)
Australia – United Kingdom	+0.032 (95% CI > 0)	1.019 (1.017 – 1.020) [No push]	β = −0.002 (−0.016 – 0.013)

How to read the columns.

- **Part 1: are the links there?** It is the mean similarity between the places of the starting country and those of the same-language country, minus the mean similarity with the controls. It shows the greater similarity defined by the test. On its own it does not separate language and geography. "95% CI > 0" means that the 95% *confidence interval*, the zone of plausible values, lies entirely above zero.
- **Part 2: does the map carry them with it?** Q says how much the estimated centre of the starting country moves closer, in proportion, to the same-language country, relative to the controls. The "placebo" values of the controls (France, Italy, Germany) lie between 0.976 and 1.035. A *placebo* is a control made where the effect sought should not be there: here, countries without the language link examined. It shows how much the estimates shift even without that link. "Push" applies the full rule: the interval of Q entirely below 1, and Q below all the placebos. By the rule fixed before the calculation, the push in the first two pairs makes the AI's translation fall.
- **Post-hoc check.** *Post-hoc* means "done afterwards". The column is explained in the point "The check after the calculation", below.

Errata. The pre-registered rules stated that the probe's pull towards the centre of the data would make the push harder to find. So a possible "falls" outcome would have been even stronger. This reassurance was not justified.

The check after the calculation.

- It was written after seeing the outcome and registered before running it. A first version proved unusable in the dry run on fake data, and it was abandoned.
- The shift of the estimated centre is decomposed, in the latitude-longitude plane, along two directions: towards the centre of the data (24.0° N, 18.6° E) and towards the same-language country (coefficient β).
- The overall reading, fixed beforehand, gives a "mixed" outcome.
- However, two directions in the plane describe any shift. It is as on a map: every shift can be stated by how far one goes north and how far east.
- The shift of Brazil is described equally well by a pull towards the centre that is stronger in latitude (0.46) than in longitude (0.29), without any contribution from Portugal.
- The check, therefore, does not attribute that component to language.

The limits.

- The check came after the outcome, and its thresholds are chosen without external references.
- The probe's coefficient comes from the map, whose places are part of the test.
- The confidence intervals resample places and estimates already obtained, without retraining the probe. They do not guarantee 95% coverage.
- Only one model, one layer and one run were used.
- The similarity of *toponyms*, that is, of place names, counts here as a language link.

How the work was done

This section gathers the procedure. In the archived text it corresponds to Chapter 5 ("Who did what"), to Appendix E, to the line on Astra's verification in Appendix F and to the contents of the package.

Who did what.

- **The author, Ugo Lanciano.** The ideas of MM presented here are the author's: the room of language, the grandmother and the aunt, the world as projection, the "irrational" method, the distinction between self-consciousness and sentient self-consciousness. The author is not a physicist. In the text the author's ideas are set out in clear words, as the author asked. The author's exact words, with the dates, are in the file AUTHOR_BACKGROUND.
- **Drafting and experiments: Claude.** An artificial intelligence, Claude, prepared at the author's request the historical reconstruction, the applied readings, the programs, the measurements and the drafting of the text. Claude also checked the sources one by one (file REFERENCE_CHECK).
- **Layout: Argon.** The form of the text (boxes, source legend, table of the test) was revised by the author with the help of another artificial intelligence, Argon. Claude then checked the content, comparing it with the previous version, and applied the corrections.
- **Verification: Astra.** Another artificial intelligence, Astra, verified the measurements and the text. It is a check between artificial intelligences, not an independent review in the strict sense. And it is *static*: Astra reads rules, programs and numbers, but does not redo the calculations.

The registration of the measurements.

- For the tests on the small network (Appendix A) and for the measurement of the space neurons (Appendix B), the original rules and programs were registered before being run. Registered means: with a cryptographic digital fingerprint (*hash*) and a time, in a local register. The fingerprint is a code calculated from the content of a file: if the file changes even slightly, the fingerprint changes.
- Corrections that emerged later are in separate files (*errata*). The registered rules were not touched.
- Some comparisons are analyses added afterwards, and they are flagged.
- The map (Appendix C) is a descriptive figure, calculated afterwards and redone once.
- The register is local transparency documentation, not an independent certification.

Astra's verifications.

1. On the tests of the small network: "ready with corrections" (10 findings, all accepted).
2. On the measurement of the space neurons and on the map: "ready with corrections" (12 findings, all accepted).
3. On the main text, nine rounds of review:
4. in the first round, "not ready" (19 findings: 18 accepted, one in part);
5. in the second, "ready with corrections" (7 findings, accepted);
6. in the third, on the corrections and on the new part on consciousness, "ready with corrections" (4 findings, accepted);
7. in the fourth, on the lightened version and on the paraphrases of the premises, "ready with corrections" (18 findings, accepted);
8. in the fifth, "ready with corrections" (7 findings: 5 accepted, 2 resolved by showing the document that was missing);
9. in the sixth, on the laid-out version, "ready with corrections" (8 findings, accepted);
10. in the seventh, on the corrections of the sixth, "ready";
11. in the eighth, on §3.2 as enlarged at the author's request, "ready with corrections" (5 findings, accepted);
12. in the ninth, on the corrections, "ready".
13. On the §3.2 test (Appendix F): outcome confirmed with corrections (12 findings, all accepted).

All the reports and replies are archived in the folder REVIEWS.

The §3.2 test and the sources.

- For the §3.2 test (Appendix F), rules and program were registered before the calculation. They were tested beforehand on fake data and run only once. The check made afterwards was also registered before being run.
- The dry run on fake data exercises the measurements and the decision. It does not exercise the loading of the data, the exclusions and the probe.
- Astra's verification of the test is static. It does not certify the reproduction of the intervals or the identity of the data.
- The bibliographic sources were checked by Claude on 6.11.2026, in four ways: on the text of the work, on

bibliographic metadata, on secondary sources (when the original could not be reached) and, for one entry only, on search results. The file REFERENCE_CHECK states entry by entry what was read.
- The data licences were verified directly on the dataset pages and in the licence files.

The versions of the archived text. The original archived text is in Italian (LEGGIMI_IT.md). The same text in English, translated by the AI (Claude) and checked against the Italian by Astra (REVIEWS), is in README.md. The composed versions for reading are HAL_9000_canta_Daisy_IT.pdf and HAL_9000_sings_Daisy_EN.pdf.

What is in the package.

- README.md: the text in English.
- LEGGIMI_IT.md: the text in Italian, the original.
- HAL_9000_canta_Daisy_IT.pdf and HAL_9000_sings_Daisy_EN.pdf: the composed versions for reading.
- AUTHOR_BACKGROUND.md and AUTHOR_BACKGROUND_IT.md: the author's sentences, with exact words, dates and context.
- FIGURE: Figure 1 in English and in Italian (PNG and SVG).
- MISURE: the rules registered before the calculations, the errata, the programs, the outcomes with the notes, the data and the local register of fingerprints and times (see MISURE/LEGGIMI_MISURE.md).
- REVIEWS: the reports of Astra (Codex), with Claude's responses.
- REFERENCE_CHECK.md: the check of the sources.
- LICENSE.md and IMPRONTE.sha256 (SHA-256 fingerprints of the files).

This version. This student version was written with an AI assistant (Claude) from the archived text in Italian, which is the author's original. It was rewritten for a school student: shorter sentences, technical words made clear, a few more examples, marked as ours. The numbers were checked with a program against the archived text. The archived text is the reference.

This English text is a translation of the Italian student version, which the author read and approved.

Glossary

The words of the framework (MM)

- **Difference (or distinction).** "this is not that": in MM it is at the base of everything.
- **GCD.** The greatest common divisor.
- **Irrational (method).** In the author's vocabulary, a method that uses distinct elements without a definition of their own, that is, without a rule formulated by a person. It does not mean illogical.
- **Local phase.** The world we measure, with places, order and distances. It arises when the differences are counted from a point.
- **Matrix.** Reality counted from a point, made of places, instants and order. For the author, language is the matrix.
- **Non-local phase.** The differences taken on their own, which have no place and no moment. It is present in every thing.
- **Non-location.** MM's own sense of "non-local": a relation that is not in a place or at an instant.
- **Placement.** Putting a thing in a place, in the local phase. In the author's example, language places the bodies of the grandmother and the aunt in two different places of its "room".
- **Premise.** A starting sentence of the author, a fixed point like a postulate.
- **Reading.** An interpretation proposed within the MM framework. Here the readings are conceptual, not laboratory results.
- **Registered.** Written and fixed with its date before the calculations, so that it cannot be changed afterwards.
- **Residue.** In the image of the loop in NUMERI, the part of itself that a system never manages to read.

The words of computing and science

- **Access consciousness, phenomenal consciousness.** For Block: the information available for reasoning and acting; lived subjective experience.
- **Accuracy.** The percentage of right answers.
- **Artificial neuron.** A small unit of calculation of a network. In language models it is a numerical value that varies continuously.
- **Backpropagation.** The method for training networks with several layers: the error is sent back and the weights are corrected.
- **Confidence interval (CI).** The zone of plausible values for a result.
- **Correlation.** A number between −1 and +1 that says how much two quantities go together.

- **Deep learning.** Learning with deep networks, that is, with many layers.
- **Dropout.** A training technique that switches off a part of the neurons at random.
- **Falsifiable.** Said of a statement that the data can refute.
- **Language model.** A neural network trained on texts, which learns to predict the next word.
- **Layer.** A row of neurons in a network.
- **Linear probe.** A simple formula that adds up the internal numbers of a network, each with its own weight, to read a piece of information.
- **Median.** The value in the middle: half of the cases below, half above.
- **Neural network.** A program made of many small linked units of calculation, which learns from examples.
- **Placebo.** A control made where the effect sought should not be there, to see how much the results shift even without that effect.
- **R^2.** A number that says how well a formula predicts a quantity: $R^2 = 1$ means perfect prediction.
- **Regularisation.** A brake put on a probe, so that it does not give too much weight to a single number.
- **Training.** Teaching a network with many examples, correcting the weights a little at each error.
- **Transformer.** A type of network from 2017, based on the mechanism of "attention".
- **Weight.** The number that says how much a link between two neurons counts.

References

Checked on 6 November 2026, entry by entry, on the text of the work, on bibliographic metadata, on secondary sources or (for one entry) on search results: the file REFERENCE_CHECK specifies for each reference what was read.

History

- McCulloch, W. S., Pitts, W. (1943). A logical calculus of the ideas immanent in nervous activity. Bulletin of Mathematical Biophysics 5(4), 115–133. doi:10.1007/BF02478259
- von Neumann, J. (1945). First Draft of a Report on the EDVAC, 30 June 1945. Reprint edited by M. D. Godfrey, IEEE Annals of the History of Computing 15(4), 27–75 (1993). doi:10.1109/85.238389
- Hebb, D. O. (1949). The Organization of Behavior. New York, Wiley.
- Turing, A. M. (1950). Computing machinery and intelligence. Mind 59(236), 433–460. doi:10.1093/mind/LIX.236.433
- McCarthy, J., Minsky, M. L., Rochester, N., Shannon, C. E. (1955). A Proposal for the Dartmouth Summer Research Project on Artificial Intelligence, 31 August 1955.
- Rosenblatt, F. (1958). The perceptron. Psychological Review 65(6), 386–408. doi:10.1037/h0042519
- Minsky, M., Papert, S. (1969). Perceptrons. MIT Press.
- Searle, J. R. (1980). Minds, brains, and programs. Behavioral and Brain Sciences 3(3), 417–424. doi:10.1017/S0140525X00005756
- Hopfield, J. J. (1982). Neural networks and physical systems with emergent collective computational abilities. PNAS 79(8), 2554–2558. doi:10.1073/pnas.79.8.2554
- Rumelhart, D. E., Hinton, G. E., Williams, R. J. (1986). Learning representations by back-propagating errors. Nature 323, 533–536. doi:10.1038/323533a0
- Smolensky, P. (1988). On the proper treatment of connectionism. Behavioral and Brain Sciences 11(1), 1–23. doi:10.1017/S0140525X00052432
- Fodor, J. A., Pylyshyn, Z. W. (1988). Connectionism and cognitive architecture: a critical analysis. Cognition 28(1–2), 3–71. doi:10.1016/0010-0277(88)90031-5
- LeCun, Y. et al. (1989). Backpropagation applied to handwritten zip code recognition. Neural Computation 1(4), 541–551. doi:10.1162/neco.1989.1.4.541
- LeCun, Y., Bottou, L., Bengio, Y., Haffner, P. (1998). Gradient-based learning applied to document recognition. Proceedings of the IEEE 86(11), 2278–2324. doi:10.1109/5.726791
- IBM, «Deep Blue» (ibm.com/history/deep-blue): the rematch with Kasparov, 3–11 May 1997.
- NVIDIA CUDA: announced in November 2006, public development kit in February 2007 (secondary source: «CUDA», en.wikipedia.org/wiki/CUDA).
- Hinton, G. E., Osindero, S., Teh, Y.-W. (2006). A fast learning algorithm for deep belief nets. Neural Computation 18(7), 1527–1554. doi:10.1162/neco.2006.18.7.1527
- Deng, J. et al. (2009). ImageNet: a large-scale hierarchical image database. CVPR 2009, 248–255. doi:10.1109/CVPR.2009.5206848
- Krizhevsky, A., Sutskever, I., Hinton, G. E. (2012). ImageNet classification with deep convolutional neural networks. Advances in Neural Information Processing Systems 25, 1097–1105.
- Russakovsky, O. et al. (2015). ImageNet Large Scale Visual Recognition Challenge. International Journal of Computer Vision 115, 211–252. doi:10.1007/s11263-015-0816-y
- Mikolov, T., Yih, W., Zweig, G. (2013). Linguistic regularities in continuous space word representations. NAACL-HLT 2013, 746–751.
- Zeiler, M. D., Fergus, R. (2014). Visualizing and understanding convolutional networks. ECCV 2014, 818–833.

doi:10.1007/978-3-319-10590-1_53
- Silver, D. et al. (2016). Mastering the game of Go with deep neural networks and tree search. Nature 529, 484–489. doi:10.1038/nature16961. For the match with Lee Sedol (Seoul, 9–15 March 2016): DeepMind, page «AlphaGo».
- Vaswani, A. et al. (2017). Attention is all you need. Advances in Neural Information Processing Systems 30. arXiv:1706.03762
- Brown, T. B. et al. (2020). Language models are few-shot learners. Advances in Neural Information Processing Systems 33, 1877–1901.
- ChatGPT, released by OpenAI on 30 November 2022 (secondary source: «ChatGPT», en.wikipedia.org/wiki/ChatGPT).
- Tiku, N. (2022). The Google engineer who thinks the company's AI has come to life. The Washington Post, 11 June 2022.
- Awards: ACM Turing Award 2018 (Bengio, Hinton, LeCun), announced on 27 March 2019; Nobel Prizes in Physics and in Chemistry 2024 (nobelprize.org).

Debate

- Block, N. (1995). On a confusion about a function of consciousness. Behavioral and Brain Sciences 18(2), 227–247. doi:10.1017/S0140525X00038188
- Damasio, A. (1999). The Feeling of What Happens: Body and Emotion in the Making of Consciousness. New York, Harcourt Brace.
- Harnad, S. (1990). The symbol grounding problem. Physica D 42, 335–346. doi:10.1016/0167-2789(90)90087-6
- Bender, E. M., Gebru, T., McMillan-Major, A., Shmitchell, S. (2021). On the dangers of stochastic parrots. FAccT '21, 610–623. doi:10.1145/3442188.3445922
- Liétard, B., Abdou, M., Søgaard, A. (2021). Do language models know the way to Rome? BlackboxNLP 2021, 510–517. doi:10.18653/v1/2021.blackboxnlp-1.40
- Olah, C. et al. (2020). Zoom in: an introduction to circuits. Distill 5(3). doi:10.23915/distill.00024.001
- Elhage, N. et al. (2022). Toy models of superposition. arXiv:2209.10652
- Mitchell, M., Krakauer, D. C. (2023). The debate over understanding in AI's large language models. PNAS 120(13), e2215907120. doi:10.1073/pnas.2215907120
- Chalmers, D. J. (2023). Could a large language model be conscious? arXiv:2303.07103
- Butlin, P. et al. (2023). Consciousness in artificial intelligence: insights from the science of consciousness. arXiv:2308.08708
- McGrath, T. et al. (2023). The Hydra effect: emergent self-repair in language model computations. arXiv:2307.15771
- Gurnee, W., Tegmark, M. (2024). Language models represent space and time. ICLR 2024. arXiv:2310.02207
- Quian Quiroga, R., Reddy, L., Kreiman, G., Koch, C., Fried, I. (2005). Invariant visual representation by single neurons in the human brain. Nature 435, 1102–1107. doi:10.1038/nature03687
- Quian Quiroga, R., Kreiman, G., Koch, C., Fried, I. (2008). Sparse but not "grandmother-cell" coding in the medial temporal lobe. Trends in Cognitive Sciences 12(3), 87–91. doi:10.1016/j.tics.2007.12.003
- Kleene, S. C. (1938). On notation for ordinal numbers. Journal of Symbolic Logic 3(4), 150–155. doi:10.2307/2267778

Data, model, historical box and MM deposits

- Alpaydin, E., Kaynak, C. (1998). Optical Recognition of Handwritten Digits. UCI Machine Learning Repository. doi:10.24432/C50P49 (licence CC BY 4.0, from the dataset page).
- Gurnee, W., Tegmark, M., world_place.csv, github.com/wesg52/world-models, commit a572f16 (MIT licence of the repository).
- Biderman, S. et al. (2023). Pythia: a suite for analyzing large language models across training and scaling. ICML 2023, PMLR 202, 2397–2430. Model EleutherAI/pythia-410m, revision 9879c9b5 (Apache-2.0 licence).
- The Maya Long Count calendar (correlations 584283 and 584285) and the inscriptions of Tortuguero (Monument 6) and La Corona: press releases from Tulane University and the University of Texas at Austin, 2012.
- Joint PASCAL VOC 2012 and ILSVRC 2012 workshop, 12 October 2012, ECCV 2012, Florence (eccv2012.unifi.it); NIPS 2012, 3–8 December 2012, Lake Tahoe (Nevada) (nips.cc).
- «Daisy Bell» (Harry Dacre, 1892) and the Bell Labs demonstration of 1961: secondary source, entry «Daisy Bell», en.wikipedia.org/wiki/Daisy_Bell (the discrepancies between sources are noted in the reference check).
- Clarke, A. C. (1968). 2001: A Space Odyssey (novel); Kubrick, S. (director) (1968). 2001: A Space Odyssey (film).
- Lanciano, U. (2026). NUMERI. La matematica del reale. ISBN 979-8187966196.
- MM 2.68.1, DOI 10.5281/zenodo.23672394.
- MM 2.72, Not at a distance, but nowhere, DOI 10.5281/zenodo.23740381.